



Practical Assignment 5 - Distributed Systems

Kubernetes and Raft



Authors:

Fernando Pastor (897113)
Peralta
Javier Murillo Jimenez (897154)

Submission date: 11/27/2025

Group: Pair: 1-6

Table of contents

Introduction:	3
Design of our solution:	3
Adaptation of the Raft server to Kubernetes	3
Execution of each node type in its specific	4
Structure of the system deployed in Kubernetes	4
Docker containers	5
Kubernetes manifests	5
Validation tests	6
Appendix	7
A. Dockerfiles	7
a. Dockerfile.srvraft.....	7
b. Dockerfile.srvraft.....	7
B. Yaml.....	8
a. cltraft_go.yaml (Client.....	8
b. srvraft_go.yaml (Headless service and StatefulSet).....	8
C. Execution and commands used.....	10
a. Image building and uploading	10
b. System startup.....	11
c. Testing	11
d. Delete everything.....	14
e. Other commands.....	14
D. The output of our script	15

Introduction:

The objective of this exercise is to deploy the distributed key/value storage service based on the Raft protocol designed in the previous exercise within a Kubernetes cluster, ensuring that each replica runs on a different worker node and that, in the event of a replica or an entire node in the cluster going down, Kubernetes automatically restarts it and Raft is able to restore its previous state without any data loss.

To do this, we had to adapt the code from the previous practice to allow connection via DNS in the Kubernetes service in a container environment.

Design of our solution:

Adapting the Raft server to Kubernetes

To adapt the code from the previous exercise, we made the necessary changes so that the nodes connect correctly to the DNS service we created (in `svraft_go.yaml`), explicitly setting the DNS address of each of the raft servers on the client, and on the raft servers, checking that the given parameter is an expected server name, constructing its address and that of the others, knowing in advance the DNS subdomain and the number of nodes in the system.

Next, we compiled the server and client .go code and obtained their executables in the `raft/` folder.

Next, we created a Dockerfile to create the images for each one, and made them accessible in a local Docker container registry (`localhost:5001`) so that they could be obtained by Kubernetes (the registry and kind were created using the provided shell file "`kind-with-registry.sh`"). In these Dockerfiles, we mounted the Alpine Linux distribution that will run on Docker and decided on the files where the container will start and work.

Shell

```
#build creates the docker image and pushes it into the local
registry ----Raft server image---

docker build -t localhost:5001/srvraft -f Dockerfile.srvraft .

docker push localhost:5001/srvraft
```

```
#---Raft client image---

docker build -t localhost:5001/clraft -f Dockerfile.clraft .

docker push localhost:5001/clraft
```

The Dockerfiles and .yaml will be available in the appendix.

Running each node type in its specific Pod

This second part can be seen reflected in the .yaml files, where we will apply these configuration manifests for Kubernetes. For the client (clraft_go.yaml), we have decided to use a basic Pod since it does not require fault tolerance. We also connect directly as clients, automatically taking the client image from the local registry. In contrast, in the Raft replicas (srraft_go.yaml), we have used a StatefulSet since we need a stable and unique identity for each node, a fixed DNS (service on port 6000), and fault tolerance. The following applies directly launch the different pods:

Shell

```
kubectl apply -f srraft_go.yaml

kubectl apply -f clraft_go.yaml
```

Structure of the system deployed in Kubernetes

Our solution is organized into three components:

- 3 Raft replicas of the storage server, executed using a StatefulSet.
- A headless Service for DNS.
- A pod that runs the client.

The StatefulSet provides us with stable DNS names for each replica, which is necessary for each node to build the RPC addresses of the others. The headless Service allows Kubernetes to register each Pod in the DNS without the need to associate a virtual IP.

Docker containers

To isolate the executables and allow their deployment in Kubernetes, we have generated two independent Docker images:

- Image of the Raft server, built from Alpine Linux and containing only the "srraft" executable.
- Image of the automatic client, also based on Alpine, which executes the "clraft" binary.

Both images, found in the appendix, have been statically built from Golang code (compiled with "CGO_ENABLED = 0") and uploaded to the local registry "localhost:5001" that Kubernetes uses as a container repository. This ensures that the cluster nodes can download the images without the need for external repositories.

Kubernetes manifests

We have designed the following manifests:

1. Client pod

Runs the "clraft" binary with a "Never" policy so that Kubernetes does not recreate it if it dies.

2. Headless service

Define the "raft-service" service, which is necessary for Kubernetes to register Raft replicas in its internal DNS. We set "clusterIP: None" so that Kubernetes does not assign a virtual IP and instead each one gets its own stable DNS name.

3. StatefulSet of replicas

The StatefulSet creates three replicas of the Raft server "srraft," each running in a different container. Using the environment variable "MINOMBREPOD," each node obtains its name within the cluster and, from there, calculates the addresses of the other replicas.

The manifests are included in the appendix.

Validation tests

For the validation tests, we created a bash script that checks whether the client logs indicate that the operations have been successful. The script performs the following tests:

1. Deployment and normal operation: after cleaning the environment, we start up the servers and client. The client performs four reads and four writes (one of them on the same key to confirm that the value is changed correctly). If everything goes well, the client prints "SUCCESS" in one of its strings. When this happens, the bash script detects it and passes the first test.
2. We kill the Pod of the leader raft node: after analyzing the logs to identify the leader node, we forcibly remove it to verify that Kubernetes regenerates it and that Raft is able to choose a new leader and continue to operate correctly with the client.
3. Node fault tolerance test (Docker crash): In bash, we identify the node (Docker container) running the new leader and stop it (simulating a network outage) to verify that the system continues to function correctly with the remaining majority. We then restore the node.

Appendix

A. Dockerfiles

a. Dockerfile.clraft

None

```
#We mount the Alpine Linux distribution in our image.

FROM alpine

#From now on, everything you do will go to /usr/local/bin.

WORKDIR /usr/local/bin

#The application inside the container listens on port 6000.

EXPOSE 6000

#Copy the Raft server executable. COPY

srvraft /usr/local/bin/srvraft #Start the

server.

ENTRYPOINT ["/usr/local/bin/srvraft"]
```

b. Dockerfile.srvraft

None

```
#Mount the Alpine Linux distribution in our image.

FROM alpine

#From now on, everything that is done will go to /usr/local/bin.

WORKDIR /usr/local/bin

#Copy the Raft server executable. COPY

clraft /usr/local/bin/clraft #Start the

server.

ENTRYPOINT ["/usr/local/bin/clraft"]
```

B. Yaml

a. cltraft_go.yaml (Client pod)

None

```
#Automatic client, connected to the Raft

service apiVersion: v1

kind: Pod

metadata:

  name: cltraft

spec:

  restartPolicy: Never

  containers:

    - name: cltraft

      image: localhost:5001/cltraft:latest

      ports:

        - containerPort: 7000
```

b. srvraft_go.yaml (Headless service and StatefulSet)

None

```
#Headless service for Kubernetes to register the pods in the DNS. apiVersion: v1

kind: Service

metadata:

  name: raft-service

  labels:

    app: raft

spec:

  clusterIP: None

  selector: #must match the label defined in the Pods.
```

```
    app: raft

    ports:
      - port: 6000
        name: raft-port
        protocol: TCP
      targetPort: 6000

    ---
    #StatefulSet of Raft replicas

    kind: StatefulSet
    apiVersion: apps/v1
    metadata:
      name: raft
    spec:
      serviceName: raft-service
      replicas: 3
      podManagementPolicy: Parallel
      selector:
        matchLabels:
          app: raft          # must correspond to
        .spec.template.metadata.labels

      template:
        metadata:
          labels:
            app: raft
        spec:
          terminationGracePeriodSeconds: 10
          containers:
            - name: srvraft
```

```

image: localhost:5001/srvraft:latest

env:
  - name: MYPODNAME      #first replica raft-0, second raft-1, etc...
    valueFrom:
      fieldRef:
        fieldPath: metadata.name
  command:
    - /usr/local/bin/srvraft
    - $(MYPODNAME)
  ports:
    - containerPort: 6000

```

C. Execution and commands used

a. Image building and uploading

Shell

```

# Compile the client and server executables

CGO_ENABLED=0 go build -o cltraft pkg/cltraft/main.go
CGO_ENABLED=0 go build -o srvraft cmd/srvraft/main.go

# Upload the images to kind-registry

----Raft server image----

docker build -t localhost:5001/srvraft -f Dockerfile.srvraft .
docker push localhost:5001/srvraft

----Raft client image---

docker build -t localhost:5001/cltraft -f Dockerfile.cltraft .
docker push localhost:5001/cltraft

```

b. System startup

Shell

```
# Apply the Kubernetes configuration manifests kubectl apply -f  
srvraft_go.yaml  
  
kubectl apply -f cltraft_go.yaml
```

c. Testing

Shell

```
bash test.bash  
  
# Next, test.bash #!/bin/bash  
  
# Configuration  
  
CLIENT="cltraft_go.yaml"  
  
SERVER="srvraft_go.yaml"  
  
POD_CLI="clientraft"  
  
# Function to run client and validate success. verify_client() {  
  
    kubectl delete pod $POD_CLI --grace-period=0 --force 2>/dev/null  
  
    kubectl apply -f $CLIENT  
  
    echo " ...Waiting for client execution..."  
  
    #The script pauses here until the pod finishes its work (Succeeded  
    or Failed status).  
  
    while true; do  
  
        #Get the current status (Running, Pending, Succeeded,  
        Failed, etc.).  
  
        PHASE=$(kubectl get pod $POD_CLI -o jsonpath='{.status.phase}'  
2>/dev/null)  
  
        # If it has already finished (successfully or unsuccessfully), we exit  
        the loop.  
  
        if [[ "$PHASE" == "Succeeded" || "$PHASE" == "Failed" ]]; then  
  
            break  
        fi  
    done
```

```
    fi

    #If it is still running, we wait a
    little. sleep 2

done

LOGS=$(kubectl logs $POD_CLI)

if echo "$LOGS" | grep -q "SUCCESS";
then echo " -> OK: Test passed.\n"
echo "-----"
kubectl get pods -o wide
echo "-----"
echo "$LOGS"
echo "-----"

else
echo " -> ERROR: The client
failed." echo "$LOGS"
exit 1

fi
}

echo "==== STARTING RAFT TEST IN KUBERNETES ==="

echo "[1] Deploying Servers..."

kubectl delete pod --all --grace-period=0 --force 2>/dev/null
kubectl apply -f $SERVER

echo " Waiting for pods to be ready..." sleep 5
echo "[2] Basic Test (Write/Read)" verify_client

# Save logs to find the leader
```

```

PREVIOUS_LOGS=$(kubectl logs $POD_CLI)

#We search for the ID of the leader to kill in the logs from the previous
execution. LEADER_ID=$(echo "$PREVIOUS_LOGS" | grep -o 'leader=[0-9]*' | tail -1
| cut -d=
-f2)

echo "[3] Deleting Leader Pod (raft-$LEADER_ID)..."
kubectl delete pod raft-$LEADER_ID --grace-period=0 --force
2>/dev/null echo " Waiting for regeneration...""

sleep 10

verify_client # Check that it is still working #

Save logs to find the leader

PREVIOUS_LOGS=$(kubectl logs $POD_CLI)

#We search for the ID of the leader to kill in the logs from the previous
execution.

LEADER_ID=$(echo "$PREVIOUS_LOGS" | grep -o 'leader=[0-9]*' | tail -1 | cut -d=
-f2)

LEADER_NODE=$(kubectl get pods -o wide | grep -E "raft-$LEADER_ID.*" | awk
'{print $7}' )

echo "[4] Stopping Docker container on node ($LEADER_ID)..."

docker stop $LEADER_ID

sleep 8

echo " Testing with node down..."

verify_client

echo "[5] Restoring node..."

docker start $LEADER_NODE

echo ""

kubectl get pods -o wide

echo "==== END ==="

```

d. Delete everything

Shell

```
kind delete cluster

docker stop kind-registry

docker rm kind-registry
```

e. Other commands

Shell

```
# Display running pods kubectl get pods -o wide

# Display Docker containers on the host system docker ps -a -s

# Display all nodes in the Kubernetes cluster kubectl get nodes -o wide

# Display logs in real time

kubectl logs -f pod_name (clientraft, raft-0, raft-1, raft-2)
```

The screenshot shows a terminal window with four distinct command outputs:

- Output 1:** \$kubectl get pods -o wide
- Output 2:** \$docker ps -a -s
- Output 3:** \$kubectl get nodes -o wide
- Output 4:** \$kubectl logs -f clientraft

The first output lists four pods: clientraft, raft-0, raft-1, and raft-2, each with a status of "Running". The second output lists several Docker containers, mostly named after the pods. The third output lists four nodes: kind-worker3, kind-worker, kind-worker2, and kind-worker4, all in a "Ready" state. The fourth output shows log entries for the clientraft pod, indicating successful writes and reads to memory locations like idx=17, 19, 21, and 23.

This screenshot shows a detailed log output for the clientraft pod. It includes a header for "INICIO DE PRUEBAS DE CLIENTE" and several log entries:

- Log entry 1: >> Escritura OK: idx=17, mandato=94, lider=1
- Log entry 2: >> Lectura OK: 'valor_inicial_1' == 'valor_inicial_1' (Esperado)
- Log entry 3: >> Escritura OK: idx=19, mandato=94, lider=1
- Log entry 4: >> Lectura OK: 'valor_inicial_2' == 'valor_inicial_2' (Esperado)
- Log entry 5: >> Escritura OK: idx=21, mandato=94, lider=1
- Log entry 6: >> Lectura OK: 'valor_inicial_3' == 'valor_inicial_3' (Esperado)

Below the log entries, there is a section titled "PRUEBA DE SOBRESCRITURA" with more log entries:

- >> Escritura OK: idx=23, mandato=94, lider=1
- >> Lectura OK: 'VALOR_ACTUALIZADO' == 'VALOR_ACTUALIZADO' (Esperado)

The log concludes with a note: "Ver logs del cliente: kubectl logs clientraft".

D. The output of our script

Shell

```
==== START RAFT TEST IN KUBERNETES ====

[1] Deploying servers...

pod "clientraft" force deleted from default namespace
pod "raft-0" force deleted from default namespace
pod "raft-1" force deleted from default namespace
pod "raft-2" force deleted from default namespace
service/raft-service unchanged
statefulset.apps/raft unchanged

Waiting for pods to be ready...

[2] Basic Test (Write/Read)

pod/clientraft created
...Waiting for client execution...
-> OK: Test passed.

-----
NAME          READY   STATUS    RESTARTS   AGE      IP           NODE
NOMINATED NODE  READINESS GATES
clientraft     0/1     Completed  0          2m18s   10.244.4.5
kind-worker2   <none>
raft-0         1/1     Running   0          2m24s   10.244.1.6
kind-worker    <none>
raft-1         1/1     Running   0          2m24s   10.244.3.3
kind-worker3   <none>
raft-2         1/1     Running   0          2m24s   10.244.2.14
kind-worker4   <none>

-----
--- START OF CLIENT TESTS ---

-> Write OK: idx=1, command=38, leader=1
```

```
-> Read OK: 'initial_value_1' == 'initial_value_1' (Expected)

-> Write OK: idx=3, command=38, leader=1

-> Read OK: 'initial_value_2' == 'initial_value_2' (Expected)

-> Write OK: idx=5, command=38, leader=1

-> Read OK: 'initial_value_3' == 'initial_value_3' (Expected)
```

OVERWRITE TEST

```
-> Write OK: idx=7, command=38, leader=1

-> Read OK: 'UPDATED_VALUE' == 'UPDATED_VALUE' (Expected)
```

--- ALL TESTS COMPLETED SUCCESSFULLY ---

[3] Deleting Leader Pod (raft-1)...

```
pod "raft-1" force deleted from default namespace

Waiting for regeneration...

pod "clientraft" force deleted from default namespace

pod/clientraft created

...Waiting for client execution...

-> OK: Test passed.\n
```

```
NAME          READY   STATUS      RESTARTS   AGE     IP           NODE
NOMINATED NODE   READINESS GATES

clientraft    0/1     Completed    0          2m17s   10.244.3.4
kind-worker3   <none>            <none>

raft-0        1/1     Running     0          4m51s   10.244.1.6
kind-worker   <none>            <none>

raft-1        1/1     Running     0          2m27s   10.244.4.6
kind-worker2   <none>            <none>
```

```
raft-2      1/1      Running     0          4m51s   10.244.2.14
kind-worker4    <none>      <none>

-----
--- START OF CLIENT TESTING ---

-> Write OK: idx=9, command=93, leader=2

-> Read OK: 'initial_value_1' == 'initial_value_1' (Expected)

-> Write OK: idx=11, command=93, leader=2

-> Read OK: 'initial_value_2' == 'initial_value_2' (Expected)

-> Write OK: idx=13, command=93, leader=2

-> Read OK: 'initial_value_3' == 'initial_value_3' (Expected)
```

OVERWRITE TEST

```
-> Write OK: idx=15, command=93, leader=2

-> Read OK: 'UPDATED_VALUE' == 'UPDATED_VALUE' (Expected)
```

--- ALL TESTS COMPLETED SUCCESSFULLY ---

```
[4] Stopping node Docker container (kind-worker4)...

kind-worker4

Testing with node down...

pod "clientraft" force deleted from default namespace

pod/clientraft created

...Waiting for client execution...

-> OK: Test passed.\n
```

```
NAME        READY     STATUS    RESTARTS   AGE       IP           NODE
NOMINATED NODE  READINESS GATES

clientraft  0/1      Completed  0          5s       10.244.3.5
kind-worker3    <none>      <none>
```

```
raft-0      1/1      Running      0          5m4s      10.244.1.6
kind-worker    <none>           <none>
```

```
raft-1      1/1      Running      0          2m40s     10.244.4.6
kind-worker2  <none>           <none>
```

```
raft-2      1/1      Running      0          5m4s      10.244.2.14
kind-worker4  <none>           <none>
```

```
-----  
--- START OF CLIENT TESTS ---
```

```
-> Write OK: idx=17, command=94, leader=1  
  
-> Read OK: 'initial_value_1' == 'initial_value_1' (Expected)  
  
-> Write OK: idx=19, command=94, leader=1  
  
-> Read OK: 'initial_value_2' == 'initial_value_2' (Expected)  
  
-> Write OK: idx=21, command=94, leader=1  
  
-> Read OK: 'initial_value_3' == 'initial_value_3' (Expected)
```

```
OVERWRITE TEST
```

```
-> Write OK: idx=23, command=94, leader=1  
  
-> Read OK: 'UPDATED_VALUE' == 'UPDATED_VALUE' (Expected)
```

```
--- ALL TESTS COMPLETED SUCCESSFULLY -----  
-----
```

```
[5] Restoring node...
```

```
kind-worker4
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
NOMINATED NODE	READINESS	GATES				
clientraft	0/1	Completed	0	5s	10.244.3.5	
kind-worker3	<none>		<none>			
raft-0	1/1	Running	0	5m4s	10.244.1.6	
kind-worker	<none>		<none>			

```
raft-1      1/1      Running     0          2m40s    10.244.4.6
kind-worker2 <none>           <none>
raft-2      1/1      Running     0          5m4s     10.244.2.14
kind-worker4 <none>           <none>

==== END ===
```